

Operating Systems Engineering


Lecture 2: From Source to Boot

Michael Engel (michael.engel@uni-bamberg.de)

Lehrstuhl für Praktische Informatik, insb. Systemnahe Programmierung

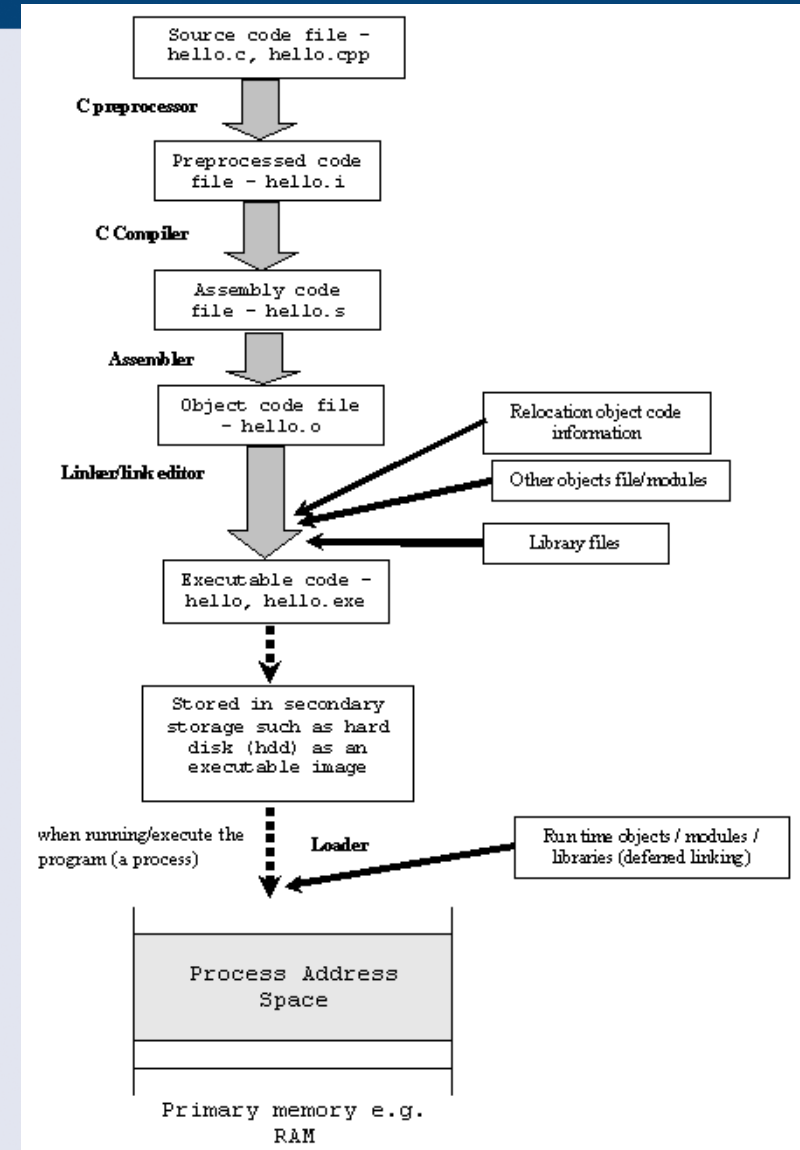
<https://www.uni-bamberg.de/sysnap>

Licensed under CC BY-SA 4.0
unless noted otherwise



- The compilation process
- Compiler, assembler, linker
- The ELF file format
- What comes before main?
- Tools
- Bare metal programming
- Writing and running the "Hello, world" example

- Preprocessor
 - Expands #includes and macros
- Compiler
 - Generates assembler source code from C source code
- Assembler
 - Generates object code from assembler source code
- Linker
 - Combines (one or) multiple object files (+ libraries) to an executable file
- Loader
 - Loads executable file into main memory



Example: From .c to .o

```
#include <stdio.h>

static void display(int i, int *ptr);

int main(void)
{
    int x = 5;
    int *xptr = &x;
    printf("In main() program:\n");
    printf("x value is %d and is stored at address %p.\n", x, &x);
    printf("xptr pointer points to addr %p which holds a value %d.\n", xptr, *xptr);
    display(x, xptr);
    return 0;
}

void display(int y, int *yptr)
{
    char var[7] = "ABCDEF";
    printf("In display() function:\n");
    printf("y value is %d and is stored at address %p.\n", y, &y);
    printf("yptr pointer points to addr %p which holds a value %d.\n", yptr, *yptr);
}
```

Example: From .c to .o (2)

```
$ gcc -c testprog1.c
```

**compiler: translates C source code
file to object file testprog.o**

```
$ ls -l
```

```
total 8
```

```
-rw-r--r-- 1 me me 629 Mar 22 13:15 testprog1.c
```

```
-rw-r--r-- 1 me me 1412 Mar 22 13:16 testprog1.o
```

**check the type of the
object file with file**

```
$ file testprog1.o
```

```
testprog1.o: ELF 32-bit LSB relocatable, intel 80386, version 1 (SYSV), not  
stripped
```

**look at the contents of the
object file with hexdump**

```
$ hexdump -C testprog1.o
```

00000000	7f 45 4c 46 01 01 01 00	00 00 00 00 00 00 00 00	.ELF.....
00000010	01 00 03 00 01 00 00 00	00 00 00 00 00 00 00 00
00000020	84 02 00 00 00 00 00 00	34 00 00 00 00 00 28 004....(
00000030	0b 00 08 00 8d 4c 24 04	83 e4 f0 ff 71 fc 55 89L\$.....q.U.

```
[...]
```

hexadecimal byte values

printable ASCII char's

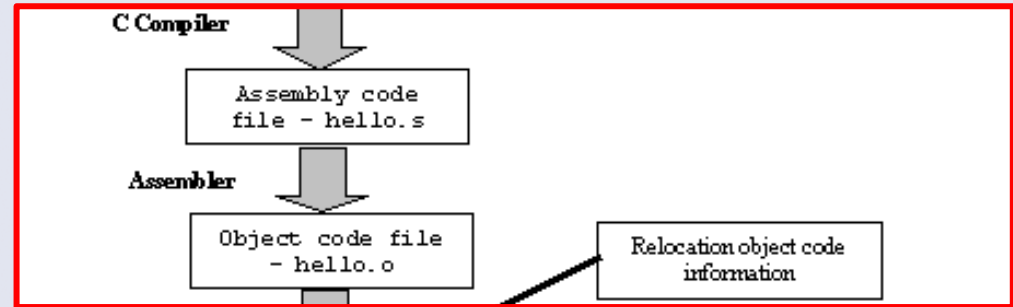
Wait a moment...

Where is our assembler source?

Only generated internally as temporary file!

Compiler option "-S" explicitly generates a .s file:

```
$ gcc -S testprog1.c
$ ls -l
testprog1.c
testprog1.o
testprog1.s
```



```
.file "testprog1.c"
.section .rodata
.LC0:
.string "In main() program:"
.align 4
.LC1:
.string "x value is %d and is stored at address %p.\n"
.align 4
.LC2:
.string "xptr pointer points to address %p which holds a value of %d.\n"
.text
.globl main
.type main, @function
main:
    leal    4(%esp), %ecx
    andl    $-16, %esp
    pushl   -4(%ecx)
    pushl   %ebp
    movl    %esp, %ebp
    pushl   %ecx
    subl    $36, %esp
    movl    $5, -12(%ebp)
    leal    -12(%ebp), %eax
```

ELF internals (1)

```
$ file testprog1.o
```

```
testprog1.o: ELF 32-bit LSB relocatable, intel 80386, version 1 (SYSV),  
not stripped
```

ELF:
Executable and Linking Format

```
$ man elf
```

```
[...]
```

The ELF header is described by the type `Elf32_Ehdr` or `Elf64_Ehdr`:

```
#define EI_NIDENT 16  
typedef struct {  
    unsigned char e_ident[EI_NIDENT];  
    uint16_t      e_type;  
    uint16_t      e_machine;  
    uint32_t      e_version;  
    ElfN_Addr     e_entry;  
    ElfN_Off      e_phoff;  
    ElfN_Off      e_shoff;  
    uint32_t      e_flags;  
    uint16_t      e_ehsize;  
    uint16_t      e_phentsize;  
    uint16_t      e_phnum;  
    uint16_t      e_shentsize;  
    uint16_t      e_shnum;  
    uint16_t      e_shstrndx;  
} ElfN_Ehdr;
```

ELF internals (2)

```
$ hexdump -C testprog1.o
```

```
00000000  7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00 |.ELF.....|
00000010  01 00 03 00 01 00 00 00 00 00 00 00 00 00 00 00 |.....|
00000020  84 02 00 00 00 00 00 00 34 00 00 00 00 00 28 00 |.....4....(.|
00000030  0b 00 08 00 8d 4c 24 04 83 e4 f0 ff 71 fc 55 89 |.....L$.....q.U.|
[...]
```

```
00000000  7f 45 4c 46 ; ELF „Magic“
00000004  01
00000005  01
00000006  01
00000007  00 00 00 00 00 00 00 00 00 00
```

/usr/include/elf.h:

```
#define ELFMAG      "\177ELF" // \177: octal
#define EI_CLASS    4 /* File class byte index */
#define ELFCLASSNONE 0 /* Invalid class */
#define ELFCLASS32   1 /* 32-bit objects */
#define ELFCLASS64   2 /* 64-bit objects */
#define ELFCLASSNUM  3
```

```
#define EI_NIDENT 16
typedef struct {
    unsigned char e_ident[EI_NIDENT];
    uint16_t      e_type;
    uint16_t      e_machine;
    uint32_t      e_version;
    ElfN_Addr     e_entry;
    ElfN_Off      e_phoff;
    ElfN_Off      e_shoff;
    uint32_t      e_flags;
    uint16_t      e_ehsize;
    uint16_t      e_phentsize;
    uint16_t      e_phnum;
    uint16_t      e_shentsize;
    uint16_t      e_shnum;
    uint16_t      e_shstrndx;
} ElfN_Ehdr;
```



`/usr/include/elf.h:`

```
#define EI_DATA          5      /* Data encoding byte index */
#define ELFDATANONE      0      /* Invalid data encoding */
#define ELFDATA2LSB      1      /* 2's complement, little endian */
#define ELFDATA2MSB      2      /* 2's complement, big endian */
```

Endianness or byte order

- Order in which the bytes of a data type using more than 8 bit are stored in memory
- Example:

Decimal number 123456789 = hexadecimal **0x075BCD15** Inspiration: Swift's "Gulliver's Travels" [4]



Little endian: least significant byte first
(at lowest address)

00000000 **15**
00000001 **CD**
00000002 **5B**
00000003 **07**

Big Endian: most significant byte first:

00000000 **07**
00000001 **5B**
00000002 **CD**
00000003 **15**

```
$ readelf testprogl.o
```

```
Usage: readelf <option(s)> elf-file(s)
```

```
Display information about the contents of ELF format files
```

```
Options are:
```

```
-a --all
```

```
Equivalent to: -h -l -S -s -r -d -V -A -I
```

```
-h --file-header
```

```
Display the ELF file header
```

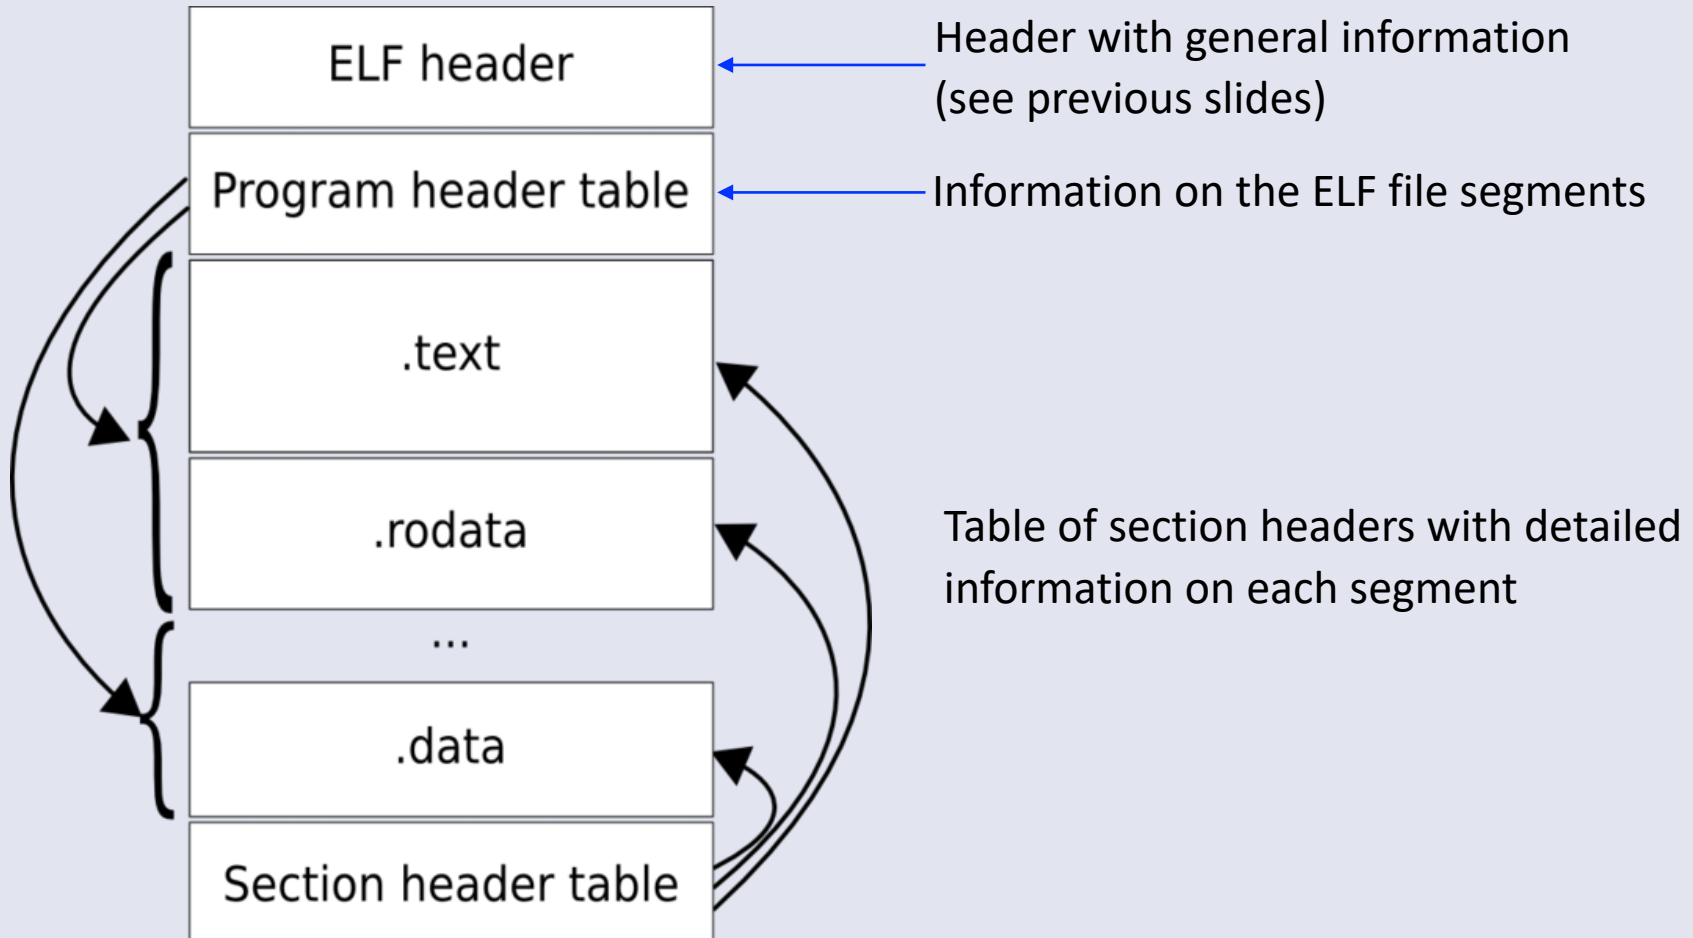
```
[...]
```

```
$ readelf -h testprogl.o
```

ELF Header:

```
Magic:  7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
Class:                                     ELF32
Data:                                       2's complement, little endian
Version:                                  1 (current)
OS/ABI:                                    UNIX - System V
ABI Version:                              0
Type:                                      REL (Relocatable file)
Machine:                                  Intel 80386
Version:                                  0x1
Entry point address:                      0x0
Start of program headers:                 0 (bytes into file)
Start of section headers:                644 (bytes into file)
Flags:                                    0x0
Size of this header:                      52 (bytes)
Size of program headers:                 0 (bytes)
Number of program headers:                0
Size of section headers:                 40 (bytes)
Number of section headers:               11
Section header string table index:      8
```

The ELF file format is standardized (see [1] and [2])



```
$ readelf -S testprog1.o
```

There are 11 section headers, starting at offset 0x284:

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
[1]	.text	PROGBITS	00000000	000034	0000db	00	AX	0	0	4
[2]	.rel.text	REL	00000000	000524	000060	08		9	1	4
[3]	.data	PROGBITS	00000000	000110	000000	00	WA	0	0	4
[4]	.bss	NOBITS	00000000	000110	000000	00	WA	0	0	4
[5]	.rodata	PROGBITS	00000000	000110	000102	00	A	0	0	4
[6]	.comment	PROGBITS	00000000	000212	00001f	00		0	0	1
[7]	.note.GNU-stack	PROGBITS	00000000	000231	000000	00		0	0	1
[8]	.shstrtab	STRTAB	00000000	000231	000051	00		0	0	1
[9]	.symtab	SYMTAB	00000000	00043c	0000c0	10		10	9	4
[10]	.strtab	STRTAB	00000000	0004fc	000026	00		0	0	1

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings)

I (info), L (link order), G (group), x (unknown)

O (extra OS processing required) o (OS specific), p (processor specific)

\$ [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?]

\$ [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?] [?]

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
[1]	.text	PROGBITS	00000000	000034	00002c	00	AX	0	0	4
[2]	.rel.text	REL	00000000	000364	000020	08		9	1	4
[3]	.data	PROGBITS	00000000	000060	000004	00	WA	0	0	4
[4]	.bss	NOBITS	00000000	000064	000000	00	WA	0	0	4
[5]	.rodata	PROGBITS	00000000	000064	000004	00	A	0	0	4

<div> <div>?</div> <div>?</div> <div>?</div> <div>?</div> <div>?</div> <div>?</div> <div>?</div> </div>	<div> <div>?</div> <div>?</div> <div>?</div> <div>?</div> <div>?</div> <div>?</div> <div>?</div> <div>?</div> </div>
text (.text)	Machine code (instructions) and entry point (address)
read only data (.rodata)	initialized constants
read/write data (.data)	initialized variables
Base Storage Segment (.bss)	uninitialised data
Symbols (.symtab)	addresses for symbolic names

```
const int a = 42;
int b = 23;
int c;

int main(void) {
    c = a + b;
    return c;
}
```

Both .data and .rodata
hold a variable using
4 bytes each = sizeof(int)

ELF section details

Offsets in ELF .o file:

?? ? ? ? ? ? ? ? ? ? ?

There are 11 section headers, starting at offset 0xd8:

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
[1]	.text	PROGBITS	00000000	000034	00002c	00	AX	0	0	4
[2]	.rel.text	REL	00000000	000364	000020	08		9	1	4
[3]	.data	PROGBITS	00000000	000060	000004	00	WA	0	0	4
[4]	.bss	NOBITS	00000000	000064	000000	00	WA	0	0	4
[5]	.rodata	PROGBITS	00000000	000064	000004	00	A	0	0	4

```
00000000  7f 45 4c 46 01 01 01 00  00 00 00 00 00 00 00 00 |.ELF.....|
00000010  01 00 03 00 01 00 00 00  00 00 00 00 00 00 00 00 |.....|
00000020  d8 00 00 00 00 00 00 00  34 00 00 00 00 00 28 00 |.....4.....(|
00000030  0b 00 08 00 8d 4c 24 04  83 e4 f0 ff 71 fc 55 89 |.....L$.....q.U.|
00000040  e5 51 8b 15 00 00 00 00  a1 00 00 00 00 8d 04 02 |.Q.....|
00000050  a3 00 00 00 00 a1 00 00  00 00 59 5d 8d 61 fc c3 |.....Y].a..|
00000060  17 00 00 00 2a 00 00 00  00 47 43 43 3a 20 28 44 |....*.....GCC: (D|
00000070  65 62 69 61 6e 20 34 2e  33 2e 32 2d 31 2e 31 29 |ebian 4.3.2-1.1)|
```

?? ? ? ? ? ? ? ? ? ? ?

0x60: 17 00 00 00 => 0x00000017 ?? ? ? ? ? ?

0x64: 2a 00 00 00 => 0x0000002a ?? ? ? ? ? ?

ELF sections and assembler code

In the assembler code:

```

$ ???????????????? .text
$ ???????????????? .globl main
.type main, @function
main:
    leal    4(%esp), %ecx
    andl    $-16, %esp
    pushl   -4(%ecx)
    pushl   %ebp
    movl    %esp, %ebp
    pushl   %ecx
    movl    a, %edx
    movl    b, %eax
    leal    (%edx,%eax), %eax
    movl    %eax, c
    movl    c, %eax
    popl    %ecx
    popl    %ebp
    leal    -4(%ecx), %esp
    ret
.size main, .-main
.comm c,4,4
.ident "GCC: (Debian 4.3.2-1.1) 4.3.2"
.section .note.GNU-stack,"",@progbits

.file "foo.c"
.globl a
.section .rodata
.align 4
.type a, @object
.size a, 4
a:
.long 42
.globl b
.data
.align 4
.type b, @object
.size b, 4
b:
.long 23
```

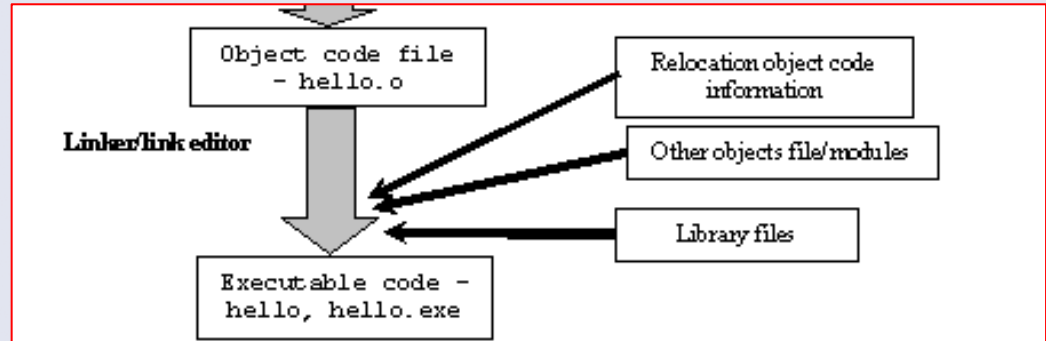
????????????

```
const int a = 42;
int b = 23;
int c;
```

```
int main(void) {
    c = a + b;
    return c;
}
```

.o object files cannot be executed directly!

- Important parts of an executable file are missing
 - crt0 – startup code
 - initialization
 - variables in .bss need to be initialized (to 0)
 - for C++: calling of constructors
 - jump to "main" function and parameter passing (argc, argv, envp)
 - libraries, e.g. libc (C standard library), have to be added
- **Addresses of variables and functions are not resolved**
 - One of the main tasks of the linker



Addresses of functions and (global) variables have to be known

- Symbol table: assigns addresses to symbolic names for functions and variables:

```
$ readelf -s foo.o
```

- Addresses are set to 0 in the object file

Symbol table '.symtab' contains 12 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	00000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	00000000	0	FILE	LOCAL	DEFAULT	ABS	foo.o
2:	00000000	0	SECTION	LOCAL	DEFAULT	1	
3:	00000000	0	SECTION	LOCAL	DEFAULT	3	
4:	00000000	0	SECTION	LOCAL	DEFAULT	4	
5:	00000000	0	SECTION	LOCAL	DEFAULT	5	
6:	00000000	0	SECTION	LOCAL	DEFAULT	7	
7:	00000000	0	SECTION	LOCAL	DEFAULT	6	
8:	00000000	4	OBJECT	GLOBAL	DEFAULT	5	a
9:	00000000	4	OBJECT	GLOBAL	DEFAULT	3	b
10:	00000000	44	FUNC	GLOBAL	DEFAULT	1	main
11:	00000004	4	OBJECT	GLOBAL	DEFAULT	COM	c

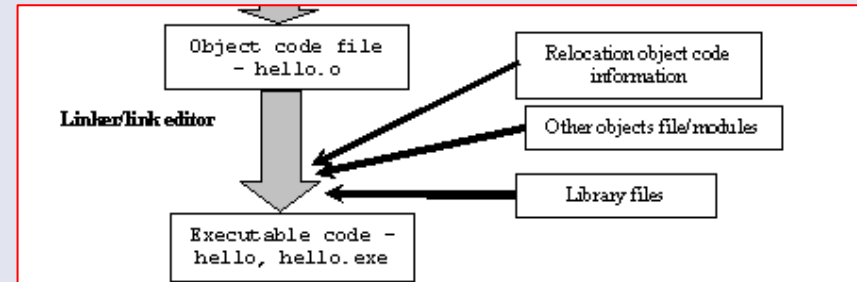


Linker receives a list of all object files required to build a program

- **.o** files (generated by the compiler or assembler): *object files*
 - can also be generated by C++, Fortran, ... compilers!
- **.a** files: "*archives*" of object files
 - *static libraries*
- **.so** files: "*shared object*"s
 - *dynamic libraries* (Windows: DLL "Dynamic Loadable Library")

Linker assigns text and data segments of the single .o files to parts of the address space for loading by the OS

- It resolves *references* to *symbols* (variables, functions)
- Controlled by a *linker script* (configuration file)



Symbol table of the linked program

- Also contains symbols of all linked libraries – a bit confusing...

```
$ gcc -o foo foo.o
```

```
$ readelf -s foo
```

Symbol table '.symtab' contains:

Num:	Value	Size	Type
------	-------	------	------

[...]

64:	08048460	4	OBJECT	GLOBAL	DEFAULT	15	a
-----	----------	---	--------	--------	---------	----	---

[...]


69:	0804956c	4	OBJECT	GLOBAL	DEFAULT	23	b
-----	----------	---	--------	--------	---------	----	---

[...]

73:	08048374	44	FUNC	GLOBAL	DEFAULT	13	main
-----	----------	----	------	--------	---------	----	------

74:	08048254	0	FUNC	GLOBAL	DEFAULT	11	_init
-----	----------	---	------	--------	---------	----	-------

75:	08049578	4	OBJECT	GLOBAL	DEFAULT	24	c
-----	----------	---	--------	--------	---------	----	---



Linker and symbol resolution (2)

Symbol table of the linked program

- Alternative: use "nm" ("names") from GNU binutils

```
$ gcc -o foo foo.o
```

```
$ nm foo
```

```
[...]
```

```
08048254 T _init
```

```
080482c0 T _start
```

```
08048460 R a
```

```
0804956c D b
```

```
08049578 B c
```

```
[...]
```

```
08048374 T main
```

?	????
T	.text
R	.rodata
D	.data
B	.bss

??????	????????
text (.text)	Machine code (instructions) and entry point (address)
read only data (.rodata)	initialized constants
read/write data (.data)	initialized variables
Base Storage Segment (.bss)	uninitialised data
??????	Symbols (.symtab)
	addresses for symbolic names

```
$ readelf -s foo
```

```
Symbol table '.symtab' contains 76 entries:
```

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
[...]							
64:	08048460	4	OBJECT	GLOBAL	DEFAULT	15	a
69:	0804956c	4	OBJECT	GLOBAL	DEFAULT	23	b
73:	08048374	44	FUNC	GLOBAL	DEFAULT	13	main
74:	08048254	0	FUNC	GLOBAL	DEFAULT	11	_init
75:	08049578	4	OBJECT	GLOBAL	DEFAULT	24	c



Linker and symbol resolution: .o files


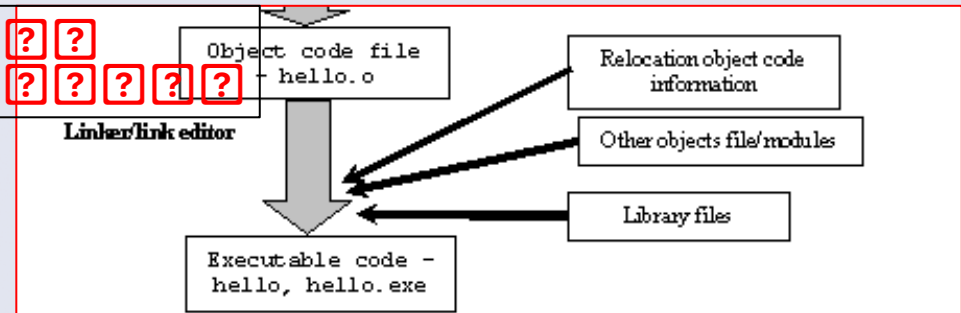
```
$ objdump -d foo.o
```

```
foo.o:      file format elf32-i386
```

```
Disassembly of section .text:
```

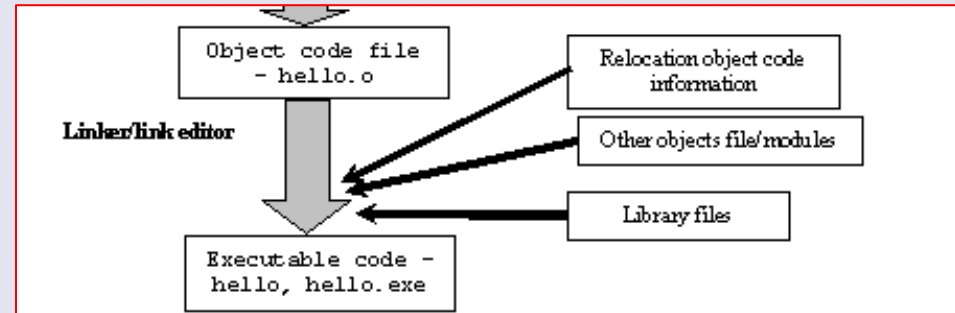
```
00000000 <main>:
```

0:	8d 4c 24 04	lea	0x4(%esp),%ecx
4:	83 e4 f0	and	\$0xffffffff0,%esp
7:	ff 71 fc	pushl	-0x4(%ecx)
a:	55	push	%ebp
b:	89 e5	mov	%esp,%ebp
d:	51	push	%ecx
e:	8b 15 00 00 00 00	mov	0x0,%edx
14:	a1 00 00 00 00	mov	0x0,%eax
19:	8d 04 02	lea	(%edx,%eax,1),%eax
1c:	a3 00 00 00 00	mov	%eax,0x0
21:	a1 00 00 00 00	mov	0x0,%eax
26:	59	pop	%ecx
27:	5d	pop	%ebp
28:	8d 61 fc	lea	-0x4(%ecx),%esp
2b:	c3	ret	



Relocation table

- Contains information about the *relative location* of the related address in the text section for addresses initialized to "0" in the object file (*relocation offset*) and the address length (R_386_32 = 32 bit)



```
$ readelf -r foo.o
```

Relocation section '.rel.text' at offset 0x364 contains 4 entries:

Offset	Info	Type	Sym.Value	Sym. Name
00000010	00000801	R_386_32	00000000	a
00000015	00000901	R_386_32	00000000	b
0000001d	00000b01	R_386_32	00000004	c
00000022	00000b01	R_386_32	00000004	c

Linker and relocations: resolution

Relocation section '.rel.text' at offset 0x364 contains 4 entries:

Offset	Info	Type	Sym.Value	Sym. Name
00000000	00000801	R_386_32	00000000	a
00000004	00000901	R_386_32	00000000	b
0000001d	00000b01	R_386_32	00000004	c
00000022	00000b01	R_386_32	00000004	c

00000000 <main>:

```

0: 8d 4c 24 04    lea    0x4(%esp),%ecx
4: 83 e4 f0      and    $0xfffffffff0,%esp
7: ff 71 fc      pushl  -0x4(%ecx)
a: 55            push   %ebp
b: 89 e5         mov    %esp,%ebp
d: 51            push   %ecx
e: 8b 15 00 00 00 00 mov    0x0,%edx
14: a1 00 00 00 00 mov    0x0,%eax
19: 8d 04 02      lea    (%edx,%eax,1),%eax
1c: a3 00 00 00 00 mov    %eax,0x0
21: a1 00 00 00 00 mov    0x0,%eax
26: 59            pop     %ecx
27: 5d            pop     %ebp
28: 8d 61 fc      lea    -0x4(%ecx),%esp
2b: c3            ret
    
```

QUESTION

ret

24

- When running an application under an operating system, the *loader* takes care of interpreting the ELF file and copying its contents into memory as required
 - But... **we are writing the operating system** :)
- What does a processor do after power on (or reset)?
 - Fetch the very first instruction from a predefined address (called reset vector)
 - The processor starts up in the most privileged mode
- A regular PC has *firmware* in non-volatile memory (flash ROM)
 - BIOS or UEFI
- RISC-V is different
 - Regular: OpenSBI and u-boot – but we will start without firmware!

- How can we convert the ELF file to something the processor can execute?
 - The ELF headers and other meta information would confuse the processor
 - qemu can actually also load ELF files, but the real hardware cannot...
 - All information that is not executable code or data has to be removed
- Solution: the **objcopy** tool:

```
riscv64-unknown-elf-objcopy -O binary hello hello.bin
```

- **Compiler**
 - Translates C (or C++, ...) source code into object files
 - We can use the GNU C compiler (GCC) or clang+LLVM
 - However, we need a **cross compiler** capable of generating instructions for a RISC-V processor!
- **Linker**
 - Links all object files of a program into a single executable
 - Part of GNU binutils (ld) – also a **cross-linker**
- **Make and Makefile**
 - Makefiles define rules to build a more complex target from multiple pieces
- **Loader**
 - Either some firmware (e.g. OpenSBI) or... no firmware at all
- **Emulator**
 - qemu – open source emulator for many different processor architectures
- **Debugger**
 - Connect to emulator and analyze code and data at runtime

- The *memory map* of a processor contains different *regions*
 - Read-only (ROM) and read-write (RAM) memory
 - I/O devices
- For our RISC-V emulated in qemu, the memory map looks like shown on the right
- Important for now:
 - UART at **0x1000_0000**
 - RAM at **0x8000_0000**
- qemu loads our hello.bin into memory at **0x8000_0000**

0xFFFF_FFFF

unallocated

0x87FF_FFFF

RAM (128 MB)

0x8000_0000

unallocated

...other I/O...

0x1000_1000

UART

0x1000_0000

Internal RISC-V I/O

0x0000_0000

BIOS ROM

It prints "Hello, world!" – what else?

```
#include <stdio.h>

int main(void) {
    printf("Hello, world!\n");
    return 0;
}
```

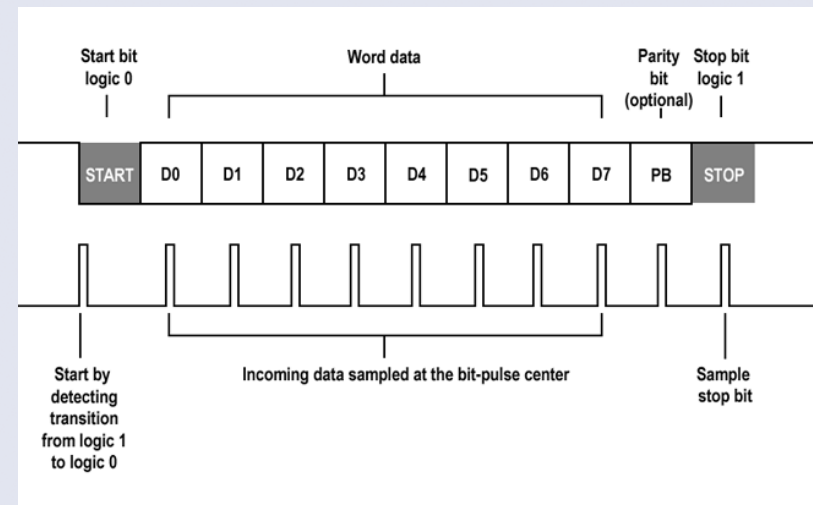
- **Problems running on bare metal:**
 - No libc – so no printf and no stdio.h header file
 - Startup code required to prepare for executing main
 - No operating system – no device drivers to actually output text

"Hello, world!" – hello.c version 2

```
void printstring(char *s) {  
    ...  
}  
  
int main(void) {  
    printstring("Hello, world!\n");  
    return 0;  
}
```

- We have to write our own function to output text
 - Here called "printstring", this is similar to `puts`, since `printf` has a much more complex parameter handling
 - How can we implement this function?

- Often there are multiple options (hardware interfaces)
 - Video output (frame buffer and display e.g. via HDMI or VGA)
 - Small LCD screen (connected via e.g. SPI or I2C)
 - Serial line (UART – universal asynchronous receiver/transmitter) – **easiest way!**
- UARTs are ICs that transmit characters in **bit serial mode**
 - One bit after the other over a **three wire connection** (transmit, receive and ground signal line)
- For real UARTs, you need to configure parameters like speed, transmit and receive parameters (e.g. number of stop bits) etc.
- We can ignore this in qemu – real hardware requires this setup!



- The 16550 UART (used in the IBM PC since 1981 as 8250) is one of the standard UART "chips" available
 - Today, it is no longer a separate IC, but a part of the system-on-chip (SoC) silicon chip
- The data sheet [5] contains a list of the chip registers and a description of their functions
- Registers are often accessible as consecutive bytes starting at the UART base address
- Registers with DLAB=1 are only accessible when LCR bit 7 = 1

Offset	DLAB	R/W	Function
0	0	R	Receive buffer reg. RBR
0	0	W	Transmit hold reg. THR
0	1	R/W	Divisor latch (LSB)
1	0	R/W	Interrupt enable reg. IER
1	1	R/W	Divisor latch (MSB)
2	X	R	Interrupt identific. reg. IIR
2	X	W	FIFO control reg. FCR
3	X	R/W	Line control reg. LCR
4	X	R/W	Modem control reg. MCR
5	X	R/W	Line status reg. LSR
6	X	R/W	Modem status reg. MSR
7	X	R/W	Scratch register

- Registers are accessible as consecutive bytes starting at the UART base address
 - For qemu: UART base = 0x1000_0000
- So RBR is at address 0x1000_0000, IER at 0x1000_0001 etc.
- Read and write operations to an address can access different functionality: read 0x1000_0000=RBR, write 0x1000_0000=THR
- Most registers combine bits for different functions
- Example: Line status register LSR
 - Transmit/Receive status plus error conditions
- Important for us now:
 - LSR at 0x1000_0005
 - Bit 5 "THR empty"

LSR bit	Function
0	Data available
1	Overrun error
2	Parity error
3	Framing error
4	Break signal received
5	THR empty
6	THR empty & line idle
7	Erroneous data in FIFO

- Input/output (I/O) devices can be run in two modes
 - **Interrupt-driven:** processor receives a signal (*interrupt*) as soon as data arrives or can be sent using the device
 - **Polling:** processor has to check continually (in a loop) if the I/O device is ready for the next piece of data
- How can our software "talk" to the UART?
 - *Registers* of the UART are available at predefined addresses
 - Reading/writing the registers result in actions of the UART
- Writing a character to a UART in polling mode looks like this (pseudocode):

```
while not (uart_transmitter_ready) { wait... }  
write character to transmit register
```

We define a structure type that describes the registers:

```
typedef char uint8_t; // defined for convenience
```

```
struct uart {  
    uint8_t THR; // transmit hold register (offset 0)  
    uint8_t IER;  
    uint8_t IIR;  
    uint8_t LCR;  
    uint8_t MCR;  
    uint8_t LSR; // line status register (offset 5)  
};
```

```
volatile struct uart* uart0 = (volatile struct uart*)0x10000000;
```

- The **uart** struct maps the registers to memory addresses
- The start address of the struct is **initialized** to 0x1000_0000;
 - The first element (THR) is at relative offset 0 (0x1000_0000)
 - ...and the LSR at relative offset 5 (0x1000_0005)

Polling the UART transmitter

```
struct uart {  
    ... // as before  
};  
  
volatile struct uart* uart0 = (volatile struct uart*)0x10000000;  
  
void putchar(char c) {  
    while ((uart0->LSR & (1<<5)) == 0)  
        ; // do nothing - wait until bit 5 of LSR = 1  
    uart0->THR = c; // then write the character  
}  
  
void printstring(char *s) {  
    while (*s) { // as long as the character is not null  
        putchar(*s); // output the character  
        s++; // and progress to the next character  
    }  
}
```


Why is the "volatile" qualifier required here?

```
volatile struct uart* uart0 =  
    (volatile struct uart*)0x10000000;
```

- Compilers are very eager to optimize away "unused" code
 - Bytes written to THR register are never read by the program
 - Bytes read from LSR are never written by the program
- The compiler cannot know that the addresses of our uart struct are not simple memory (RAM), but have *side effects* – controlling the UART
- Thus, volatile tells the compiler not to optimize away accesses to these variables
 - Volatile is one of the most misunderstood features of C [6]

- A few assembler instructions are required to get to main
 - Initialize the **stack pointer SP**
 - Jump to main
 - Finally, an endless loop if main returns

```
        .section .text
        .global _entry
_entry:
        la    sp, stack0    // stack0 is a 4096 byte array
        li    a0, 4096      //
        add   sp, sp, a0    // set sp to end of stack

        jal   main          // subroutine call to main
loop:
        j     loop          // endless loop if main returns
```

So we need a stack in hello.c

```
__attribute__((aligned (16))) char stack0[4096];  
  
struct uart {  
    ...  
};  
  
// ...and the rest...
```

- The stack is simply a region in memory, represented by a C array
 - Stores return addresses, local variables, etc.
 - Required for function calls in C
- Here, stack0 is simply a 4096 byte array
 - Aligned to a multiple of 16 bytes (required by RISC-V)
- boot.S sets the stack pointer to the first byte **after** that array
 - This works since the stack pointer is first decremented (-4) before data is written to the stack!

- The **linker script** defines where in memory the different parts of the executable program are to be located ("stolen" from xv6...)

```
OUTPUT_ARCH( "riscv" )
ENTRY( _entry )
SECTIONS
{
    /* ensure _entry is at 0x80000000
     * where qemu -kernel jumps.
     */
    . = 0x80000000;

    .text : {
        *(.text .text.*)
        . = ALIGN(0x1000);
        PROVIDE(etext = .);
    }
    /* continued on the right side */
```

```
    .rodata : {
        . = ALIGN(16);
        *(.rodata .rodata.*)
    }

    .data : {
        . = ALIGN(16);
        *(.data .data.*)
    }

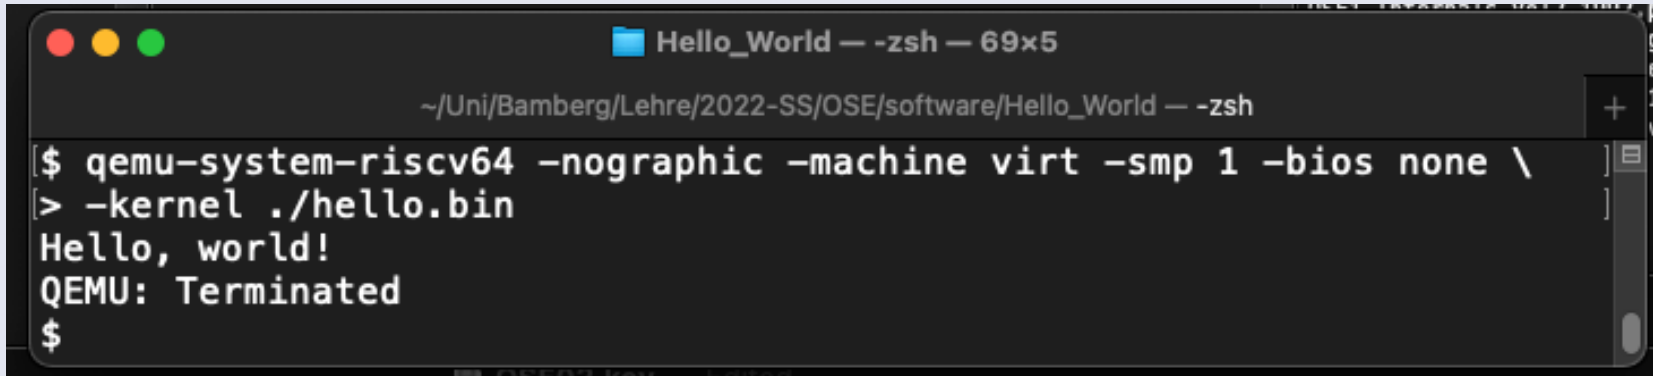
    .bss : {
        . = ALIGN(16);
        *(.bss .bss.*)
    }

    PROVIDE(end = .);
}
```


- Compile the C and the assembler source file
- Link both into an ELF executable "hello"
- Convert the ELF file into a binary `hello.bin` so that `qemu` can execute it
 - Here, as a shell script – we'll build a makefile next time!
- **-ffreestanding**: don't depend on `libc` (or `crt0`) – we have our own

```
#!/bin/bash
riscv64-unknown-elf-gcc -g -mcmodel=medany -mno-relax \
    -c boot.S
riscv64-unknown-elf-gcc -g -mcmodel=medany -mno-relax \
    -c hello.c
riscv64-unknown-elf-gcc -g -ffreestanding -fno-common \
    -nostdlib -mno-relax -mcmodel=medany \
    -Wl,-T hello.ld *.o -o hello
riscv64-unknown-elf-objcopy -O binary hello hello.bin
```

- Run in qemu... **YES!**



```

Hello_World — -zsh — 69x5
~/Uni/Bamberg/Lehre/2022-SS/OSE/software/Hello_World — -zsh
$ qemu-system-riscv64 -nographic -machine virt -smp 1 -bios none \
> -kernel ./hello.bin
Hello, world!
QEMU: Terminated
$
```

- Parameters to qemu:
 - **-nographic**: run without video output, UART output in terminal
 - **-machine virt**: specific machine type (with virtual I/O)
 - **-smp 1**: emulate a single-processor system
 - **-bios none**: run without a bios
 - **-kernel ./hello.bin**: use our hello.bin program as kernel, mapped to 0x8000_0000
- Exit qemu with the following combination: first Control-A, then press X

Some regs have different names for R and W at the same offsets

- This can be handled by using C **unions** – two names for one element
- The correct use is the responsibility of the programmer (i.e., you can still write to RBR, but it would nevertheless end up in the THR register)

```
struct uart {  
    union {  
        uint8_t THR; // W = transmit hold register (offset 0)  
        uint8_t RBR; // R = receive buffer register (also offset 0)  
        uint8_t DLL; // RW = divisor latch low (off. 0 when DLAB=1)  
    };  
    union {  
        uint8_t IER; // RW = interrupt enable reg. (offset 1)  
        uint8_t DLH; // divisor latch high (offset 1 when DLAB=1)  
    };  
    union {  
        uint8_t IIR; // R = interrupt identif. reg. (offset 2)  
        uint8_t FCR; // W = FIFO control reg. (also offset 2)  
    };  
    // ...rest as before  
};
```

- **We made it!**
 - Our first program runs bare metal on a RISC-V (emulator)
- Such a simple "Hello world" program can be useful to...
 - check if your compilation toolchain is working
 - check if your emulator works
 - check if your assumption about the interaction of hardware and software is correct
- In the next lecture, we will take a detailed look at the RISC-V processor architecture
 - ...and then we will dive into the xv6 operating system

1. System V Application Binary Interface Edition 4.1 (1997-03-18)
2. ELF TIS Committee, *Tool Interface Standard (TIS) Executable and Linking Format (ELF) Specification, Version 1.2*
Online unter <https://www.uclibc.org/docs/elf.pdf>
3. John R. Levine, *Linkers and Loaders*, MKP 1999, ISBN 1-55860-496-0 – Online (beta) at <https://www.iecc.com/linker/>
4. Jonathan Swift, *Gulliver's Travels, or Travels into Several Remote Nations of the World*, 1726
5. UART 16550 data sheet
https://www3.tuhh.de/osg/Lehre/SS21/V_BSB/doc/uart-16550d.pdf
6. Eric Eide and John Regehr, *Volatiles Are Miscompiled, and What to Do about It*, Proc. of EMSOFT 2008